

Peer-to-Peer Downloading

CS 360 Internet Programming

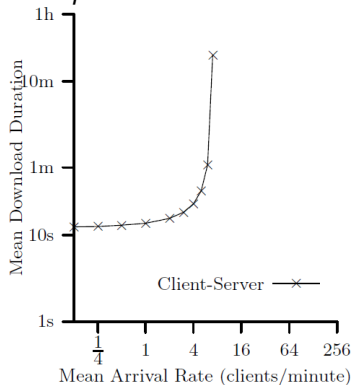
Daniel Zappala

Brigham Young University
Computer Science Department

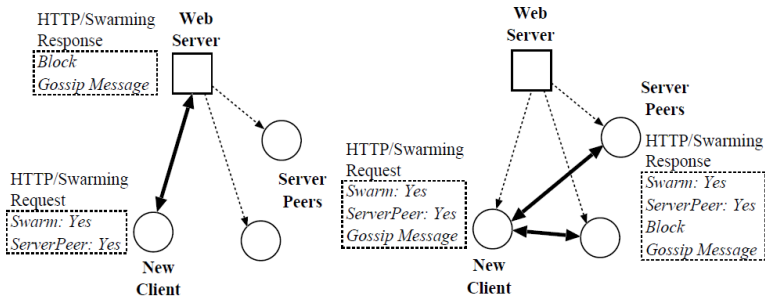
Motivation

- web servers are limited in their scalability
 - for a given web server, the more clients that need to be served, the slower they access the content
 - eventually the wait becomes so long, TCP connections time out
- CDN is a good solution, but expensive

serving a 1 MB file from a 1Mbps server



Swarming



(a) A new client gets a block and a gossip message.


(b) The client finds and downloads blocks in parallel.

- peer-to-peer content delivery
 - serve blocks of a file to some peers
 - have these peers serve this content to other peers

Scalability

- self scaling: the more clients that participate – the more bandwidth is available

would need a 28 Mbps server to handle this load



Flash Crowds

- 6 clients/minute steady state
- client-server flash crowd: 1 hour of 12 clients/minute
- swarming flash crowd: 1 hour of 120 clients/minute




Fig. 6. Client-Server flash crowd




Fig. 7. Swarming flash crowd

Basic Mechanisms

- bootstrapping
 - download a .torrent file from a web server
 - contact the listed *tracker* for a list of peers
- peer discovery
 - periodically contact tracker if current peers don't have content you need or they leave the system
- content location
 - check with each peer to determine which blocks they have
 - download rarest blocks first

Content Delivery

- seed
 - web site must run a *seed* which has the entire file so that at least one peer is available that has the whole file
 - other peers that download the file may *linger* and act as a seed during this time
- parallel download
 - keep j connections active
 - download a block at a time from each connection
- incentives
 - serve content to k connections at a time
 - try to serve the connections that give you the most content
 - periodically serve content to a random connection to see if it can do better than a current connection

The .torrent File

- **announce**: the URL of the tracker
- **info**: a dictionary with the keys
 - **name**: suggested name for the file
 - **piece length**: number of bytes in each block of the file, default 256K
 - **pieces**: list of SHA-1 hashes of the blocks of the file
 - **length**: the length of the file in bytes
- other options for multi-file torrents

Communicating with the Tracker

- client sends an HTTP GET with:
 - SHA-1 hash of the info value from the .torrent file
 - peer identifier for this peer, chosen randomly
 - IP address of the peer
 - port number of the peer
 - number of bytes uploaded and downloaded so far
 - the number of bytes left to download
 - event: started, completed, or stopped
- tracker sends a response
 - interval: number of seconds the peer should wait before contacting the tracker again
 - list of peers that are active in the system, including peer identifier, IP address, and port

Connections

- peers make many TCP connections to other peers
 - when you finish downloading a block, tell your peers
 - you may **choke** (send no data) or **unchoke** (send data) each peer
 - you may tell each peer you are **interested** in their data or **not interested**
- if you are unchoked and interested, you may request blocks from that peer
 - always download from as many peers as possible
 - request blocks in random order

Choking

- two purposes
 - limit the number of uploading connections to provide good TCP performance
 - provide incentives - I will upload to you if you upload to me
- goals
 - cap number of uploads
 - avoid rapid switches between choking and unchoking
 - reciprocate to peers that let you download
 - try unused connections periodically to improve performance
- algorithm
 - operates once every 10 seconds
 - unchokes four peers for which you have the best download rates and who are interested in you (rolling 20 second average)
 - choose a single peer who gets unchoked regardless of its upload rate and who is not currently, rotate this peer every 30 seconds ,give new connections a better chance to be that peer
- seeds unchoke those peers with the best upload rates

Interesting Issues

- how well is it working?
 - what happens when people use it?
 - how well do incentives work?
- how can I make it work better?
 - what is the right peer selection strategy?
 - what is the right block selection strategy?