

Performance Measurement and Workload Models

CS 360 Internet Programming

Daniel Zappala

Brigham Young University
Computer Science Department

Motivation: Content Creators

- goal: maximize profit
- how do people use my web site?
 - what led people to buy this book? recommendation list? ad on another site?
 - what led people to *not* buy this book? negative reviews? ordering process too complicated?
- how many people use my web site?
 - ad revenue
 - where do they come from?
 - where are *not* coming from?

Motivation: Web Hosting and ISPs

- goal: maximize profit and performance
- accounting
 - bill people for the bandwidth they use
- system performance
 - replicate content to balance load
 - evaluation of a server to find and examine bottlenecks
 - effect of server parameters, e.g. number of active threads
 - measure the effectiveness of caching: number of hits, bandwidth saved

Motivation: Network Researchers

- goal: understand the network and how well it is working
- understanding
 - determine how many hosts have deployed a new protocol or feature (e.g. how many caches use Etag for validation?)
 - measure and model Internet traffic
- protocol development: improve performance
 - led to HTTP/1.1 persistent connections
 - major role in determining TCP performance
 - evaluate cache replacement, validation, and prefetching algorithms

Passive Measurement

- examine log files
- place a monitor “inside” the network
 - promiscuous mode on Ethernet means you get all packets
 - modify host to log all IP packets to a very large disk
- advantages
 - can capture all HTTP and TCP behavior
 - can examine a detailed timeline of all events
 - non-invasive for the web server and clients
- limitations
 - hard to capture all packets at line speed
 - limited to examining what happened, not what *may* happen

Active Measurement

- generate user requests on your own to examine web server and user performance
- locating user agents
 - wide variety of connection speeds
 - wide variety of network locations
 - what distributions?
- generating requests
 - which web pages?
 - how many requests per client?
 - use distributions collected from past experiments
- collecting data
 - DNS time
 - TCP connection time
 - HTTP request time

Simulation

- model a system and examine its behavior in a simulator
- requires modeling
 - the Internet (connectivity, bandwidth, delay)
 - Internet traffic
- complexity depends on level of simulation
 - packet-level simulation: treat each packet as an event, which must be handled at each router and host in the network
 - session-level simulation: treat each file transfer as an event, which must be handled only at the hosts in the network

Simulation vs Measurement

- advantages of simulation
 - network can be tightly controlled delay, bandwidth, competing traffic
 - can often simulate at a larger scale than a measurement study
 - experiments in a simulated network run in simulated time (much faster than real time)
 - can often vary more factors for a more complete study
- advantages of measurement
 - you learn different things when you actually have to build it
 - more convincing results: it works in the “real world”
 - can use real users and real traffic

Deciding Between Simulation and Measurement

- can you implement your solution?
 - time allocated to design, spec, testing, debugging
 - is there an existing framework or library that would make your job easier?
- can you deploy your solution?
 - do you have access to enough computers or friends?
 - what operating system and libraries will you use?
 - can you use PlanetLab?
- how many factors do you want to study?

Experimental Methodology

- important components of an experimental methodology for network performance measurement
 - 1 topology
 - 2 workload
 - 3 data collection
 - 4 metrics

Topology

- what does the network look like?
 - LAN
 - campus
 - PlanetLab
 - Internet
- simulated networks must be generated
 - how many nodes?
 - how many links?
 - link and node characteristics - bandwidth, delay, queue size
 - network characteristics - organizations, ISPs

Workload

- what load will you put on your solution and therefore on the network?
- *example: web server*
 - how often do clients make a request to the web server?
 - constant arrival rate
 - exponential arrival rate
 - how many requests does a client make and what is the time between requests?
 - what size are the objects that are available on the web server?
 - what is the popularity of each of the objects?
 - what is the available bandwidth to each client?

Data Collection

- collect a log of important events in your system
- **important: collect data, not metrics**
 - if you collect metrics, you will invariably find that you need to change a metric or add another one and now you need to re-run all of your experiments
 - collecting data takes more room but if you generalize it then you can compute any metrics you want later on
- *example: web server*
 - thread time threadID start
 - thread time threadID stop
 - queue time size

Metrics

- how do you know whether your solution is solving the problem?
- how do you know it is doing it better than other solutions?
- *hint*: look at previous solutions and see what metrics they use
- *example: web server*
 - response time
 - client bandwidth
 - percentage of time each thread is busy
 - queue size over time
 - average, standard deviation, percentiles, distributions, etc.
- note: knowing whether a difference is significant requires a good statistics background

Two Approaches

- trace-driven workload
 - use a log of actual requests to a server
 - reproduces an actual workload
 - avoids statistical analysis
 - no flexibility
 - may not be representative of other situations
 - mingles load with performance (the rate at which a user can make requests depends on how fast it can download a page)
 - hard to collect, requires extensive storage
- stress testing
 - send requests as fast as possible to induce high load
 - can quickly identify performance bottlenecks
 - not realistic

Synthetic Workloads

- characterize the load placed on a system using a mathematical/statistical model
- want it to be repeatable, realistic, parameterized
- advantages
 - derived from a model that can be analyzed, validated, and compared to other models
 - can capture general properties of real traffic without being limited to a particular trace
 - can explore performance in a controlled setting by adjusting the parameters of the model

Modeling Workload

- *example: web server log - how can we characterize the size of a response message?*
- mean can be skewed by high variance
 - response messages: 4100, 4700, 4200, 20,000, 4000 bytes
 - mean size = 7400 bytes
- median doesn't represent variance
 - median size = 4200 bytes
 - response messages: 4100, 4700, 4200, 4800, 4000 bytes has same median
- variance, standard deviation provide more information but they are also relatively simple
- need probability distributions to provide a more complete picture

Probability Density Function (pdf)


- describes the likelihood for a continuous random variable to occur over an interval
- defined by a function $f(x)$ such that
 - $f(x) \geq 0$ for all values of x
 - the total area under the graph is 1

$$\int_{-\infty}^{\infty} f(x) dx = 1 \quad (1)$$

- probability of the interval $(a, b]$ is given by $\int_a^b f(x) dx$ for any two numbers a and b



Cumulative Distribution Function (CDF)

- describes the likelihood that a continuous random variable will have a value less than or equal to x
- $F(x) = P(X \leq x)$



Exponential Distribution


- $F(x) = e^{-\lambda x}$
- frequently used to model the time interval between successive random events that happen at a constant average rate
- memoryless: $\Pr(T > t + s \mid T > t) = \Pr(T > s)$
- mean $\mathbf{E}[X] = \frac{1}{\lambda}$




Pareto Distribution

- frequently used to model resource sizes
 - Italian economist Vilfredo Pareto (1848 - 1923)
 - 1906 observation: *20% of the population own 80% of the property*
- $P(X > x) = (\frac{x}{b})^{-a}, x \geq b > 0, b = \text{minimum possible value}$
 - also called a “heavy-tailed” distribution
 - linear when plotted on a log-log scale
 - Wikipedia uses $a = k, b = 1$ in graphs below

$$\text{pdf} = a \frac{b^a}{x^{a+1}}, x \geq b > 0$$



$$\text{cdf} = 1 - (\frac{b}{x})^a$$



Generating Random Variates

- want to be able to generate a random workload that follows a particular distribution
 - variate: observation or event
 - *example: generate the arrival of the next web requests using an exponential distribution*
- easy for an exponential function: inverse transform method
 - $P(X \leq x) = cdf = 1 - e^{-\lambda x}$
 - solve for x as a function of the uniform $[0,1]$ random variable U
 - $x = -\frac{1}{\lambda} \ln(1 - U)$
 - choose a random number between 0 and 1, and calculate the corresponding value of x
- Pareto
 - $x = \frac{b}{(1-U)^{\frac{1}{a}}}$

Discrete Probability

- probability mass function (pmf)
 - probability that a discrete random variable is exactly equal to some value
 - rank resources from highest to lowest probability and plot probability for each rank
- typically use Zipf's Law to model resource popularity

Zipf's Law

- in a corpus of natural language, the frequency of a word is inversely proportional to its rank
 - Harvard linguist George Kingsley Zipf
 - *example: in the Brown Corpus, "the" accounts for 7% of all words, "of" accounts for about 3.5% of all words, etc*
- $P(r) = kr^{-1}$, where r is the rank of an object and k is constant such that all probabilities sum to 1
 - more generally, $P(r) = kr^{-c}$ for some constant c
 - heavy-tailed
- similar to Pareto
 - Zipf-like distributions can be viewed as Pareto distributions if you exchange some variables
 - like a Pareto distribution, plotting ranks on a log-log scale should result in a straight line

Determining the Right Distribution

- ① determine workload parameters
- ② measure some real traffic
- ③ analyze measurement data to construct a statistical model of the workload
- ④ validate the model against additional real workloads

Generating a Web Workload

- ① generate session (user) arrivals using an exponential distribution
- ② for each session, generate the number of pages and the number of objects per page using Pareto distributions
- ③ for object size, use a Pareto distribution
- ④ for object popularity use a Zipf distribution