

Streaming Audio and Video

CS 360 Internet Programming

Daniel Zappala

Brigham Young University
Computer Science Department

Types of Streaming

- stored audio and video
 - content pre-recorded and stored at server
 - user begins playback before entire file received
 - content played continuously, at same rate as original recording
 - user can pause, rewind, fast-forward, index content
- live audio and video
 - content is sent live, rather than pre-recorded, still continuous
 - higher expectation that delays are short, no pauses unless user requests it
- real-time interactive audio and video
 - user interaction - conversations
 - highest expectation for low delay – must be below human thresholds for interactive conversation to be possible
 - delay $< 150\text{ ms}$ not perceptible, $150 - 400\text{ ms}$ acceptable, $> 400\text{ ms}$ not tolerable

Challenge

- the Internet offers only best-effort service
 - no guarantee on delay (affects live and interactive streaming)
 - no guarantee on throughput (affects video streaming)
- **multimedia applications are typically delay sensitive but loss tolerant**
- data applications are typically loss sensitive but delay tolerant

Internet Phone Example

- streaming audio: alternating talk spurts, silent periods
 - packets generated only during talk spurts
 - talk spurt = 64 kbps
 - sample every 20 ms: 160 bytes per packet
 - send packets using UDP
- ideally, during a talk spurt a packet arrives every 20ms and you play it
- **loss**: will make the stream pause or sound choppy
- **delay**: must be $< 400\text{ ms}$, $< 150\text{ ms}$ is ideal
- **jitter**: packets experience different amounts of queueing delay

Quality of Service

- performance guarantees: bandwidth, delay, jitter
- **how will the Internet support QoS?**
- **do nothing**
 - use application-level techniques to mitigate effects of delay and jitter
 - add more bandwidth as needed
 - use CDNs for stored content
- **differentiated service**
 - classify and treat some traffic differently
 - must be careful of network neutrality – treat all voice or video the same (even your competitor's), unless the user has paid for special treatment

Coping with the Internet

- applications make the best out of best-effort service
 - compression
 - client buffering
 - UDP + error correction or plain TCP
 - multiple encodings – select the best that can be delivered
 - adaptive HTTP streaming
- media player responsibilities
 - remove jitter
 - decompress stream
 - conceal errors
 - user controls: pause, rewind, fast forward

Audio Compression


- ① analog signal sampled at constant rate
 - telephone: 8,000 samples/second
 - CD music: 44,100 samples/second
 - ② each sample quantized (rounded)
 - for example 256 possible values
 - ③ each quantized value represented by bits
 - for example 8 bits for 256 values
 - 9,000 samples/second, 256 quantized values : 64,000 bps
 - ④ receiver converts digital data back to analog signal, some quality reduction
- typical rates
 - CD: 1.411 Mbps
 - MP3: 96, 128, 160 kbps
 - VoIP: 14 kbps

Video Compression

- video is a sequence of images displayed at a constant rate
 - for example 24 frames/second
- each frame is divided into an array of pixels
- each pixel is represented by some bits
- can eliminate spatial and temporal redundancy in and between frames
- example video rates
 - MPEG 1 (CD-ROM): 1.5 Mbps
 - MPEG 2 (DVD): 3 - 6 Mbps
 - MPEG 4 (Internet): < 1 Mbps
- layered video
 - hierarchical encoding into layers
 - additional layers produce increasing resolution
 - send number of layers that path to receiver can handle

Client Buffering

- use buffering to compensate for delay, jitter
 - play packets from the buffer at a certain rate
 - want to prevent draining the buffer too early
 - want to prevent delaying the packets too much: 400 ms maximum, but less is better



UDP vs TCP


• UDP

- server sends at rate appropriate for client (often ignores congestion control)
 - send rate = encoding rate = constant rate
 - fill rate = constant rate - packet loss
- short playout delay (2-5 seconds) to compensate for network delay jitter
- recover errors if time permits

• TCP

- send at maximum possible rate under TCP: congestion control
- fill rate fluctuates due to TCP congestion control
- need a larger playout delay to smooth TCP delivery rate
- HTTP/TCP passes more easily through firewalls

Fixed Playout Delay



- play packet i at $t_i + q_i$
 - t_i : time the packet was generated
 - small q : better real-time interaction
 - large q : fewer missed playouts


Adaptive Playout Delay

- goal: minimize playout delay, infrequent missed playouts
- for start of talk spurt, play packet i at $p_i = t_i + d_i + Kv_i$
 - d_i = estimate of delay (EWMA)
 - v_i = estimate of delay variation (EWMA)
 - K = constant, e.g. $K = 4$
- for other packets, $p_j = t_j + p_i - t_i$
 - use the same offset from t as beginning of talk spurt
 - application may denote start of talk spurt
- because playout delay is adaptive, silent periods may be compressed or elongated

Forward Error Correction (FEC)

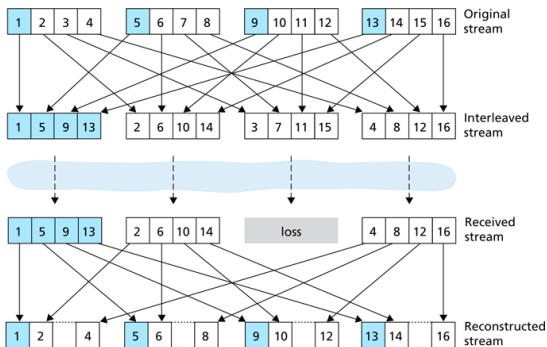
- add some redundant data to the stream
 - for every group of n packets, create a redundant packet: exclusive-OR of the n original packets
 - send $n + 1$ packets, increasing bandwidth by $1/n$
- receiver can reconstruct the stream with any n packets
 - must wait for n packets before playout
- trade-offs
 - larger n : less bandwidth
 - smaller n : shorter playout delay, smaller chance of two packets out of n being lost
- simple version of FEC: see RFC 2733

FEC With Lower Quality Redundant Data



- add lower-resolution audio to each packet
- substitute lower quality when needed: Free Phone, RAT
- many other more complex kinds of FEC available

Interleaving



- interleave smaller (5 ms) pieces among packets
- if a packet is lost, still have most of the stream
- no redundancy overhead, but added playout delay

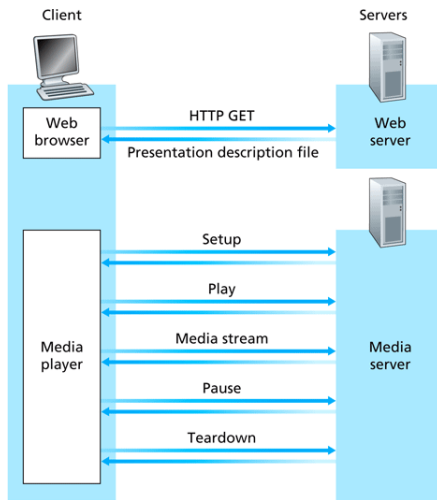
Adaptive HTTP Streaming

- used for streaming stored video (TV shows, movies)
 - Move Networks
 - Microsoft (Smooth Streaming in Silverlight)
 - Adobe (Flash)
 - Apple
- architecture
 - break video into 2 second segments
 - create multiple encodings (resolutions) of each segment
 - store segments on CDNs near users
 - video player downloads the best quality it can in 2 seconds, so it is ready to play for the next 2 second window
 - uses TCP and HTTP
- no buffering
- instantaneous switching to a new part of the movie

RTSP: User Control


- RFC 2326
- client-server, application layer protocol
- user can pause, rewind, fast forward, index, etc
- does not define
 - how media is streamed over network: packet format
 - transport protocol used
 - client buffering
- messages are *out-of-band* with respect to the streaming media: port 554

RTSP Example



RTP: Real-Time Protocol

- specifies packet structure for audio and video streams
 - payload identification
 - sequence number
 - timestamp
- used for interoperability between multimedia applications
- runs on top of UDP
- RFC 1889




RTCP: Real-Time Control Protocol

- works with RTP
- each participant in RTP session periodically transmits RTCP control packets to other participants
 - statistics from sender and receiver on session quality
- use feedback to control performance, e.g. congestion control, type of encoding
- packets sent using same mechanism as data (unicast or multicast)

RTCP Packets

- **source description:**
sender's name, host name, email address, IP of associated RTP stream
- **sender report:** IP of RTP stream, current time, number of packets sent, number of bytes sent
- **receiver report:** fraction of packets lost, last sequence number, average jitter




RTCP Extras

- synchronization of different media streams in an RTP session
 - timestamps in audio and video streams tied to sampling clocks, not wall-clock time
 - each RTCP sender report contains timestamp of most recent RTP packet and wall-clock time of when it was created
- bandwidth scaling
 - RTCP will limit control traffic to 5%
 - e.g. if sender transmits video at 2 Mbps, then RTCP tries to use 100 kbps
 - 75% of rate for receivers, 25% to sender
 - receivers try to share 75% equally
 - receiver determines RTCP transmission period by dividing average RTCP packet size by allocated rate

SIP: Session Initiation Protocol

- vision
 - all telephone and video conference calls take place over the Internet
 - people are identified by names and email addresses rather than by phone numbers
 - you can reach the person you are calling wherever she is on the Internet
- call setup
 - start and end call
 - agree on media type and encoding
- map name and email to IP address
- call management
 - add new streams during call
 - change encoding during call
 - invite other users to call
 - transfer and hold calls

SIP Call Setup



- example assumes Alice knows Bob's IP address
- Alice provides port number, IP address, preferred encoding
- Bob responds with new port number, IP address, preferred encoding
- SIP messages can use TCP or UDP, default port 5060

SIP Extras

- encoding negotiation
 - may not have requested encoder
 - may prefer a different one
 - respond with **606 Not Acceptable** and list encoders
 - sender can send a new invitation with new encoder
- can reject a call
 - *busy, gone, payment required, forbidden*
- media can use RTP or any other protocol
- syntax is similar to HTTP

Name Translation and User Location

- need to map user name or email address to IP address
 - mobility
 - changing IP addresses due to DHCP
 - many different IP devices per user
 - call forwarding
- **SIP registrar**: clients register to provide current location and IP address (similar to instant messaging)
- **SIP proxy**: find callee on behalf of caller (similar to DNS server)

Session Initiation Example

