

Thread Synchronization

CS 360 Internet Programming

Daniel Zappala

Brigham Young University
Computer Science Department

A Simple Example

```
1 void echo()  
2 {  
3     chin = getchar();  
4     chout = chin;  
5     putchar(chout);  
6 }
```

- shared method among multiple threads
- unpredictable execution sequence
 - Thread1 executes up to conclusion of input function, then interrupted: Thread1 reads 'x'
 - Thread2 executes completely through the procedure: Thread2 reads and prints 'y'
 - Thread1 starts again: Thread1 input variable has 'y'!

Concurrency Problems

- *demonstration*

Concurrency Problems

- *demonstration*
- problems
 - can't predict the speed with which threads will execute and therefore when a resource will be accessed
 - if synchronization is not used, errors will be rare but they will occur
 - errors are hard to duplicate and debug since they are nondeterministic

Mutual Exclusion

- need to protect shared resources (e.g. global variable, shared data structures) among multiple processes or threads
 - atomic access to method
 - when Thread2 tries to enter method, block it until Thread1 is finished
- may involve processes or threads interleaved in time on a single processor or running in parallel on a multiprocessor machine
- result of process or thread must be independent of the speed of execution of other concurrent processes

Mutual Exclusion

- **critical section**: shared portion of code that must be executed by one thread at a time
 - thread must mark the critical section because OS doesn't know where it is
- **starvation**: one or more threads are prevented from ever executing critical section
- **deadlock**: situation in which no thread can make progress because they are all waiting for a critical section
- must ensure data coherence, e.g. atomic access to a database

Mutual Exclusion Requirements

- only one thread has access to critical section at a time
- halting in non-critical section must not interfere with other threads
- no indefinite wait for critical section, i.e. no starvation or deadlock
- if no thread in critical section, then no wait to enter
- no assumptions about process speeds or number of processors
- thread may only spend finite time within critical section

Solutions

- software
 - assume no support from OS, hardware, or language
 - historic algorithms: Dekker, Peterson, Lamport
 - difficult to get right, to generalize
- hardware
 - disable interrupts: single processor machines
 - no other process can run until they are re-enabled
 - limits flexibility of OS to schedule threads, doesn't work for multiprocessors
 - atomic machine instructions: test-and-set
- operating system support

Test-and-Set Example

```
1  boolean test-and-set(int i) {
2      if (i == 0) {
3          i = 1;
4          return true;
5      } else {
6          return false;
7      }
8  }
9
10 int bolt = 0;
11 void method(int i) {
12     while (true) {
13         while (!test-and-set(bolt))
14             /* do nothing */
15             /* critical section */
16             bolt = 0;
17         /* remainder */
18     }
19 }
```

Test-and-Set Pros and Cons

- advantages
 - applicable to any number of processes on either a single processor or multiple processors sharing main memory
 - it is simple and therefore easy to verify
 - it can be used to support multiple critical sections: each section gets its own variable
- disadvantages
- busy-waiting consumes processor time
- starvation is possible when more than one process waits
- deadlock
 - low priority process has the critical region
 - higher priority process needs it
 - higher priority process obtains the processor and waits for the critical region

Operating System Support

- mutex and condition variable
 - mutex: lock that allows only one thread into a critical section
 - condition variable: signal conditions between threads
- semaphore
 - when one thread is in the critical section, others may **wait** by sleeping
 - when thread is done with critical section, it wakes one other thread with a **signal**
- monitor
 - programming language construct that makes it easier to declare and use a critical section
 - construct a class with methods, only one thread may access a method of the class at a time
- message passing
 - synchronization by explicitly exchanging messages
 - define a mailbox and enter critical section when a message is waiting

Mutex

- lock that allows only one thread into a critical section

```
1 #include <pthread.h>
2
3 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
4
5 int pthread_mutex_lock(pthread_mutex_t *mutex);
6 int pthread_mutex_trylock(pthread_mutex_t *mutex);
7 int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- must initialize the mutex first
- `pthread_mutex_lock()` will block if mutex is already locked
- `pthread_mutex_trylock()` will return EBUSY if mutex is locked
- *demonstration*

Busy Waiting

```
1 while running {  
2     c = NULL;  
3     pthread_mutex_lock(&mutex);  
4     if queue.not_empty() {  
5         c = queue.dequeue();  
6     }  
7     pthread_mutex_unlock(&mutex);  
8     if c {  
9         /* handle connection */  
10    }  
11 }
```

- must busy wait until a connection is available
- wastes CPU time on a server that does not handle many connections

Condition Variables

```
1 #include <pthread.h>
2 pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
3
4 int pthread_cond_wait(pthread_cond_t *cond,
5                       pthread_mutex_t *mutex);
6
7 int pthread_cond_signal(pthread_cond_t);
```

- must initialize the condition variable first
- `pthread_cond_wait()` will block until the condition is signaled; the thread now owns the mutex as well
- need a corresponding `pthread_cond_signal()` to wake up

Using Condition Variables

```
1 while running {  
2   c = NULL;  
3   pthread_mutex_lock(&mutex);  
4   while queue.empty() {  
5     pthread_cond_wait(&cond,&mutex);  
6   }  
7   c = queue.dequeue();  
8   pthread_mutex_unlock(&mutex);  
9   /* handle connection */  
10 }
```

- process inserting into queue should signal condition when queue goes from empty to having at least one item
- **must re-check queue status when conditional wait returns**
- no guarantee that queue will be empty when you return

Timed Wait and Broadcast Signals

```
1 #include <pthread.h>
2
3 int pthread_cond_timedwait(pthread_cond_t *cond,
4                             pthread_mutex_t *mutex,
5                             const struct timespec *abstime);
6
7 int pthread_cond_broadcast(pthread_cond_t *cond);
```

- `pthread_cond_timedwait()` needs an absolute time; use `clock_gettime()` and add the length of time you want to wait
- `pthread_cond_broadcast()` wakes up all threads waiting for a signal

Semaphores

- semaphore is a shared integer variable
 - initialized ≥ 0
- `wait(s)`: wait for a signal on semaphore `s`
 - decrements semaphore, blocks if < 0
 - process suspends until signal is sent
 - OS uses a queue to hold waiting processes
- `signal(s)`: transmit a signal to semaphore `s`
 - increments semaphore
 - if ≤ 0 then unblock someone
- `wait()` and `signal()` are atomic operations and cannot be interrupted

Types of Sempahores

- binary semaphore
 - only one process at a time may be in the critical section
- counting semaphore
 - a fixed number of processes > 0 may be in the critical section
- OS determines whether processes are released from queue in FIFO order or otherwise; usually FIFO in order to prevent starvation

Using Semaphores

```
1 semaphore s = 1;
2
3 void thread(int i) {
4     while (true) {
5         wait(s);
6         /* critical section */
7         signal(s);
8         /* remainder */
9     }
10 }
```

- semaphore protects critical section
- can set s to > 1 to let more than one process in the critical section
 - $s \geq 0$: number that can enter
 - $s < 0$: number that are waiting

POSIX Semaphores

```
1 #include <semaphore.h>
2
3 int sem_init(sem_t *sem, int pshared, unsigned int value);
4 int sem_wait(sem_t * sem);
5 int sem_trywait(sem_t * sem);
6 int sem_post(sem_t * sem);
```

- `sem_init()`: sets initial value of semaphore; `pshared = 0` indicates semaphore is local to the process
- `sem_wait()`: suspends process until semaphore is > 0 , then decrements semaphore
- `sem_trywait()`: returns `EAGAIN` if semaphore count is $= 0$
- `sem_post()`: increments semaphore, may cause another thread to wake from `sem_wait()`
- *demonstration*

Producer Consumer Problem

- one or more producers are generating data and placing them in a buffer
- one or more consumers are taking items out of the buffer
- only one producer or consumer may access the buffer at any time

Producer Consumer with Infinite Buffer

producer:

```

1 while (True) {
2     /* produce item v */
3     b[in] = v;
4     in++;
5 }

```

consumer:

```

1 while (True) {
2     while (in <= out)
3         /* do nothing */;
4     w = b[out];
5     out++;
6     /* consume item w */
7 }

```

Producer Consumer using Binary Sempahores

```
1 int n;  
2 sem_t s, delay;  
3 sem_init(&s,0,1);  
4 sem_init(&delay,0,0);
```

producer:

```
1 while (True) {  
2     produce();  
3     sem_wait(&s);  
4     append();  
5     n++;  
6     if (n==1) sem_post(&delay);  
7     sem_post(&s);  
8 }
```

consumer:

```
1 int m;  
2 sem_wait(&delay);  
3 while (True) {  
4     sem_wait(&s);  
5     take();  
6     n--;  
7     m = n;  
8     sem_post(&s);  
9     consume();  
10    if (m==0) sem_wait(&delay);  
11 }
```

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s?*
- 2 *What is the purpose of semaphore delay?*

Looking at the Code ...

- ① *What is the purpose of semaphore s ?*
- ② *What is the purpose of semaphore delay?*
- ③ *Why is semaphore s initialized to 1 but semaphore delay is initialized to 0?*

Looking at the Code ...

- ① *What is the purpose of semaphore s ?*
- ② *What is the purpose of semaphore delay?*
- ③ *Why is semaphore s initialized to 1 but semaphore delay is initialized to 0?*
- ④ *Why does the consumer need to wait on delay first?*

Looking at the Code ...

- ① *What is the purpose of semaphore s ?*
- ② *What is the purpose of semaphore delay?*
- ③ *Why is semaphore s initialized to 1 but semaphore delay is initialized to 0?*
- ④ *Why does the consumer need to wait on delay first?*
- ⑤ *Why does the producer signal wait only when $n == 1$?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*
- 2 *What is the purpose of semaphore delay?*
- 3 *Why is semaphore s initialized to 1 but semaphore delay is initialized to 0?*
- 4 *Why does the consumer need to wait on delay first?*
- 5 *Why does the producer signal wait only when $n == 1$?*
- 6 *Why does the consumer need a local variable m ?*

Important Insights

- two purposes for semaphores
 - **mutual exclusion**: semaphore s controls access to critical section
 - **signalling**: semaphore *delay* coordinates when the buffer is empty: consumer waits if buffer is empty, producer signals when buffer becomes non-empty
- avoid race conditions
 - m keeps a local copy of the data protected by the semaphore so that it can be accessed later
 - reduces amount of processing inside the critical section

Producer Consumer using Counting Sempahores

```
1 sem_t s, n;  
2 sem_init(&s,0,1);  
3 sem_init(&n,0,0);
```

producer:

```
1 while (True) {  
2     produce();  
3     sem_wait(&s);  
4     append();  
5     sem_post(&s);  
6     sem_post(&n);  
7 }
```

consumer:

```
1 while (True) {  
2     sem_wait(&n);  
3     sem_wait(&s);  
4     take();  
5     sem_post(&s);  
6     consume();  
7 }
```

Looking at the Code ...

- 1 *What is the purpose of semaphore s?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*
- 2 *What is the purpose of semaphore n ?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*
- 2 *What is the purpose of semaphore n ?*
- 3 *Why is semaphore s initialized to 1 but semaphore n is initialized to 0?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*
- 2 *What is the purpose of semaphore n ?*
- 3 *Why is semaphore s initialized to 1 but semaphore n is initialized to 0?*
- 4 *Why can the producer signal n every time an item is added to the buffer ?*

Looking at the Code ...

- 1 *What is the purpose of semaphore s ?*
- 2 *What is the purpose of semaphore n ?*
- 3 *Why is semaphore s initialized to 1 but semaphore n is initialized to 0?*
- 4 *Why can the producer signal n every time an item is added to the buffer ?*
- 5 *Can the producer swap the signals for n and s ?*

Looking at the Code ...

- ① *What is the purpose of semaphore s ?*
- ② *What is the purpose of semaphore n ?*
- ③ *Why is semaphore s initialized to 1 but semaphore n is initialized to 0?*
- ④ *Why can the producer signal n every time an item is added to the buffer ?*
- ⑤ *Can the producer swap the signals for n and s ?*
- ⑥ *Can the consumer swap the waits for n and s ?*

Important Insights

- more elegant solution, can't do this easily with mutexes
- n : semaphore value is number of items in buffer
 - if $n == 0$, consumer must wait
 - can swap `sem_post(&n);` and `sem_post(&s);` in producer and be OK
 - can't swap `sem_wait(&n);` and `sem_wait(&s);` in consumer: otherwise consumer enters and then waits and deadlocks the producer!
- ordering of semaphore operations is important
- would like programming language support to help with organizing mutual exclusion code: monitors

Producer Consumer Circular Buffer

```

1 sem_t s, n, e;
2 sem_init(&s,0,1);
3 sem_init(&n,0,0);
4 sem_init(&e,0,BUFFER_SIZE);

```

producer:

```

1 while (True) {
2     produce();
3     sem_wait(&e);
4     sem_wait(&s);
5     append();
6     sem_post(&s);
7     sem_post(&n);
8 }

```

consumer:

```

1 while (True) {
2     sem_wait(&n);
3     sem_wait(&s);
4     take();
5     sem_post(&s);
6     sem_post(&e);
7     consume();
8 }

```

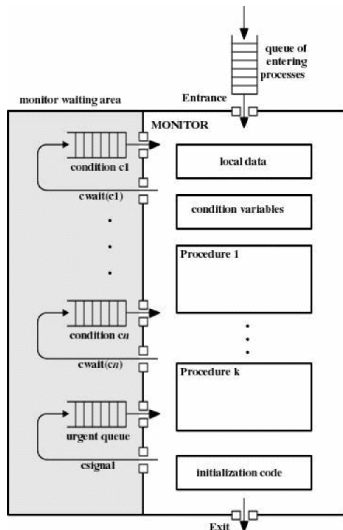
Looking at the Code ...

- 1 *What is the difference between semaphore e and semaphore n?*

Monitor

- difficult to get semaphores right
 - match wait and signal
 - put in right order
 - scattered throughout code
- **monitor**: programming language construct
 - equivalent functionality
 - easier to control
 - mutual exclusion constraints can be checked by the compiler
 - used in versions of Pascal, Modula, Mesa
 - Java also has a Monitor object but compliance cannot be checked at compile time

Hoare Monitor



- monitor can only be entered through procedures
- data can only be accessed by procedures
- only one process or thread in monitor at any time
- may suspend and wait on a condition variable
- like object-oriented programming with mutual exclusion added in

Hoare Synchronization

- `cwait(c)`: suspend on condition `c`
- `csignal(c)`: wake up one thread waiting for condition `c`
 - do nothing if no threads waiting (signal is lost)
 - different from semaphore (number of signals represented in semaphore value)

Producer Consumer with a Hoare Monitor

```
1 char buffer[BUFSIZE];
2 int in = 0, out = 0, count = 0;
3 condition notfull, notempty;
```

append:

```
1 if count == N
2     cwait(notfull);
3 buffer[in] = x;
4 in = (in + 1) % BUFSIZE;
5 count++;
6 csignal(notempty);
```

take:

```
1 if count == 0
2     cwait(notempty);
3 x = buffer[out];
4 out = (out + 1) % BUFSIZE;
5 count--;
6 csignal(notfull);
```

Producer Consumer with a Hoare Monitor

producer:

```
1 char x;  
2 while (True) {  
3     x = produce();  
4     append(x);  
5 }
```

consume:

```
1 char x;  
2 while (True) {  
3     x = take();  
4     consume(x);  
5 }
```

- advantages

- moves all synchronization code into the monitor
- monitor handles mutual exclusion
- programmer handles synchronization (buffer full or empty)
- synchronization is confined to monitor, so it is easier to check for correctness
- write a correct monitor, any thread can use it (semaphores need to be placed properly in all threads)

Lampson and Redell Monitor

- Hoare monitor requires that signaled thread must run immediately
 - thread that calls `csignal()` must exit the monitor or be suspended
 - for example, when `notempty` condition signaled, thread waiting must be activated immediately or else the condition may no longer be true when it is activated
 - usually restrict `csignal()` to be the last instruction in a method (Concurrent Pascal)
- Lampson and Redell
 - replace `csignal()` with `cnotify()`
 - `cnotify(x)` signals the condition variable, but thread may continue
 - thread at head of condition queue will run at some future time
 - must recheck the condition!
 - used in Mesa, Modula-3

Producer Consumer with a Lampson Redell Monitor

```
1 char buffer[BUFSIZE];
2 int in = 0, out = 0, count = 0;
3 condition notfull, notempty;
```

append:

```
1 while (count == N)
2     cwait(notfull);
3 buffer[in] = x;
4 in = (in + 1) % BUFSIZE;
5 count++;
6 cnotify(notempty);
```

take:

```
1 while (count == 0)
2     cwait(notempty);
3 x = buffer[out];
4 out = (out + 1) % BUFSIZE;
5 count--;
6 cnotify(notfull);
```

Lampson Redell Advantages

- allows processes in waiting queue to awaken periodically and reenter monitor, recheck condition
 - prevents starvation
- can also add `cbroadcast(x)`: wake up all processes waiting for condition
 - for example, append variable block of data, consumer consumes variable amount
 - for example, memory manager that frees k bytes, wake all to see who can go with k more bytes
- less prone to error
 - process always checks condition before doing work


What Can You Do?

- emulate a Lampson Redell Monitor with mutex and condition variables or semaphores
 - create a class with private data only
 - use the same mutex or semaphore to protect all class methods
 - use condition variables or semaphores to replace `cwait()` and `cnotify()`
 - use while loops to recheck conditions
- take your semaphores and move them inside the method call instead of outside of it (see circular buffer implementation)

Message Passing

- needed for distributed systems: no shared memory
- two primitives
 - `send(destination,message)`
 - `recv(source,message)`
- blocking
 - Operating System may block sender or receiver unless you use non-blocking calls
 - when sender returns, this does not mean the message has been delivered, just accepted by the transport protocol

Message Passing with a Mailbox



- create a mailbox abstraction (managed by a separate thread)
 - threads send messages to the central mailbox
 - address message directly to another thread
 - threads poll mailbox to get messages
- provide mutual exclusion
 - use a null message as a token
 - process that gets the token can enter the critical section
 - if token is not available, block until you get it

Producer Consumer with a Mailbox

```
1 int null = 0;
2 create_mailbox(mayproduce);
3 create_mailbox(mayconsume);
4 for (int i = 1; i <= CAPACITY; i++)
5     send(mayproduce, null);
```


producer:

```
1 message pmsg;
2 while (True) {
3     receive(mayproduce, pmsg);
4     pmsg = produce();
5     send(mayconsume, pmsg);
6 }
```

consumer:

```
1 while (true) {
2     receive(mayconsume, cmsg);
3     consume(cmsg);
4     send(mayproduce, null);
5 }
```

Direct Communication



- have a thread in charge of each shared data structure or file
- web server example
 - send a message to request a new connection
 - send a message to log statistics
 - send a message to log a request

Unix Domain Socket Address Structure

```
1 struct sockaddr_un {  
2     sa_family_t    sun_family;    // AF_LOCAL  
3     char            sun_path[108]; // path name  
4 }
```

- path name must be null terminated

UNIX Domain Server

```
1 struct sockaddr_un server;
2 char *filename = "/tmp/mysocket";
3 bzero(&server, sizeof(server));
4 server.sin_family = AF_UNIX;
5 strncpy(server.sun_path, filename, sizeof(server.sun_path) - 1);
6
7 s = socket(PF_UNIX, SOCK_STREAM, 0)
8 if (!s) {
9     perror("socket");
10    exit();
11 }
12 if (bind(s, &server, sizeof(server)) < 0) {
13     perror("bind");
14     exit();
15 }
```

- call `unlink(filename)` when finished with socket

UNIX Domain Client

```
1 struct sockaddr_un server;
2 char *filename = "/tmp/mysocket";
3 bzero(&server, sizeof(server));
4 server.sin_family = AF_UNIX;
5 strncpy(server.sun_path, filename, sizeof(server.sun_path) - 1);
6
7 s = socket(PF_UNIX, SOCK_STREAM, 0)
8 if (!s) {
9     perror("socket");
10    exit();
11 }
12 if (connect(s, &server, sizeof(server)) < 0) {
13     perror("connect");
14     exit();
15 }
```


Deadlock Definition and Conditions

- *permanent blocking of a set of processes or threads that either compete for system resources or communicate with each other*
- conditions
 - ① **mutual exclusion**: only one thread may use a resource at a time
 - ② **hold-and-wait**: a thread keeps one resource while waiting for another
 - ③ **no preemption**: a thread can't be forced to release a resource
 - ④ **circular wait**: a cycle of threads waiting for each other
- if first three conditions hold, then deadlock is possible if circular wait occurs
- depends on execution order!

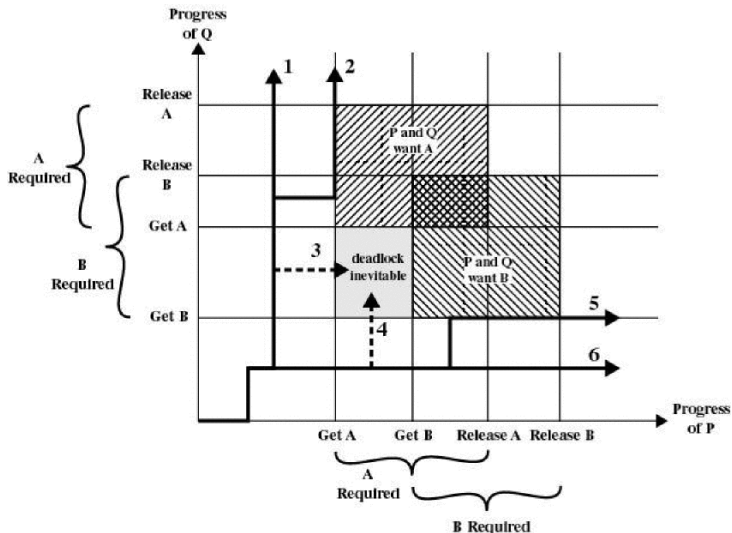
Example

```
1 Thread P
2 ...
3 Get A
4 ...
5 Get B
6 ...
7 Release A
8 ...
9 Release B
```

```
1 Thread Q
2 ...
3 Get B
4 ...
5 Get A
6 ...
7 Release B
8 ...
9 Release A
```

- *is deadlock possible?*

Deadlock Possibilities



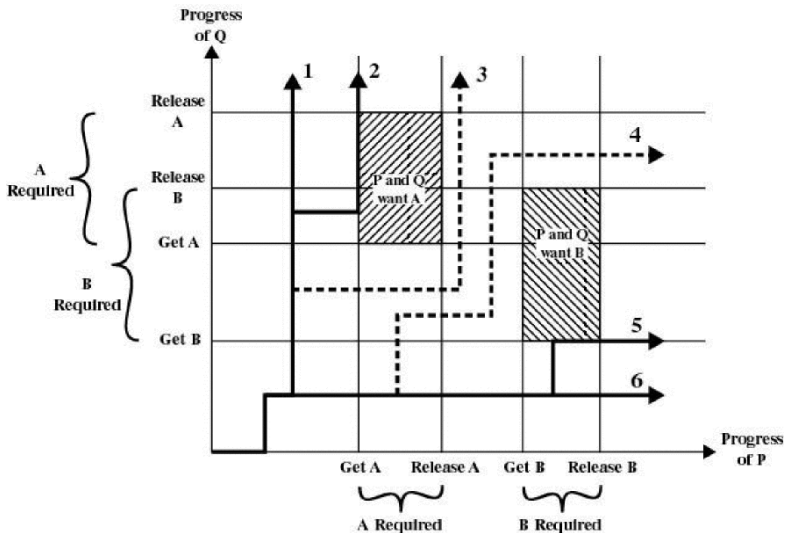
Revised Sample Code

```
1 Thread P
2 ...
3 Get A
4 ...
5 Release A
6 ...
7 Get B
8 ...
9 Release B
```

```
1 Thread Q
2 ...
3 Get B
4 ...
5 Get A
6 ...
7 Release B
8 ...
9 Release A
```

- *is deadlock possible?*

Deadlock Avoided



Simple Deadlock Prevention

- prevent one of the conditions from happening
- simplest to prevent is **hold-and-wait**: hold only one resource at a time
- can also prevent **circular wait**: impose ordering on resources