

Web Proxies and Caching

CS 360 Internet Programming


Daniel Zappala

Brigham Young University
Computer Science Department

Web Proxies

Proxy

- *an intermediary program that acts as both a server and a client for the purpose of forwarding requests*
- accepts requests from other clients and handles them either internally or by passing them on to other servers
- examples
 - **caching** responses
 - **anonymizing** requests
 - **filtering** content



Transparent Proxies

- **transparent proxy**: does not modify the request other than superficially
 - caches place identifying information in headers
- **non-transparent proxy**: may modify the request or response
 - anonymize request
 - filter content
 - compress response


Anonymizing Requests

- anonymizing proxy hides IP address of client (client is not necessarily in the same organization)
 - may not hide the User-Agent header
 - may not drop cookies
- **onion routing**: setup a sequence of proxies along an unpredictable path, using encryption at each step
 - prevents eavesdropping
 - prevents traffic analysis
 - <http://tor.eff.org/>
 - http://en.wikipedia.org/wiki/Onion_Routing

Filtering Content

- examine application-level HTTP messages to block access to certain content
 - examine URL in the GET and compare to a blacklist of web sites
 - compare URL against a list of banned keywords:
[anonymizing searches](#) often blocked
 - examine response and compare to a list of banned keywords
- BYU CS department uses DansGuardian dansguardian.org


Interception Proxies



- all web traffic is diverted to the proxy, regardless of user preference
- diverting traffic
 - ① router must examine TCP header on all packets
 - ② a TCP packet going to port 80 is diverted to the proxy
 - ③ proxy must accept packets for any destination address going to port 80
 - ④ proxy then performs its functions – caching, filtering
- breaks the rules and layering of IP, but so do firewalls
- a reality for most major campuses and organizations

Web Caching

Motivation



- problem: bandwidth bottleneck at a server
- example
 - 100 Mbps connection
 - 100,000 KB typical web page size (with embedded content)
 - 125 requests per second
- buy more bandwidth
 - 1 Gbps (10 x more bandwidth)
 - 1250 requests per second
- hard to scale

Web Caching

- may return a cached object rather than contacting the origin web server
- need a **cache consistency** protocol - check whether objects in cache are up-to-date
- need a **cache replacement algorithm** - determine which objects to save when the cache is full
- **hit rate** determined by cache replacement algorithm, workload (object popularity, object size)


Reverse Proxies

- a cache that sits in front of web server
 - provide access to a server behind a firewall
 - centralize security concerns at one server
 - balance load among a set of back-end servers
 - provide one URL space for many different web sites
- can use Apache as either forward or reverse proxy

Benefits

Caching Benefits for Users

- faster download - web cache is usually on the local network, where there is more available bandwidth
- lower latency - shorter propagation delay for closer servers
- less congestion - fewer users sharing bandwidth, local networks are usually over-provisioned



Caching Benefits for Web Servers

- lower load on the server - can handle more users
- lower cost since it uses less bandwidth
- *but some servers want all the traffic because they receive revenue for ads on the site*
- solutions: no caching on ads, pay per click-through on ad instead of per visit, survey users like with traditional media

Caching Benefits for the Network

- the network as a whole
 - less traffic traversing the Internet, since it stays on local networks
 - reduces congestion - lower delay, lower packet loss
 - improves throughput - faster transfer times
- Internet Service Providers
 - each ISP pays its upstream provider based on its access link speed (bits per second) or the actual amount of traffic sent over the link
 - big incentive to provide web caches for their users - reduces the amount of traffic on the access link, which reduces their overall cost

What is Cacheable?

What is Cacheable?

- **Expires** header
 - date after which the response is considered stale and must be revalidated
 - cache does not need to revalidate item each time it has a cache hit
- **ETag** header
 - tag specific to a resource
 - decouples cache validation from expiration times, since clocks are not synchronized
 - cache uses **If-Match** header to check if the cached item is the same
- **Vary** header
 - lists fields that may vary in responses (e.g. language)
 - cache must check that these fields are the same in the request and the cached response

Server Control over Caching

- **Cache-Control** header specifies directives that **MUST** be obeyed by a cache regardless of its own algorithms
- restrictions on what is cacheable
 - **public**: item **MAY** be cached even if normally not cacheable (e.g. responses that have an Authorize field)
 - **private**: item **MUST NOT** be cached (intended for one user)
 - **no-cache**: **MUST NOT** be returned by a cache without validation
- restrictions on what may be stored
 - **no-store**: cache **MUST NOT** store any part of the request or response

Browser Control over Caching

- expiration mechanism
 - **max-age**: maximum age client wants from cache
 - **max-stale**: gives maximum staleness client wants from cache
 - **min-fresh**: client wants a response that will still be fresh for a minimum amount of time
- cache revalidation and reload
 - **end-to-end reload**: user wants item from origin server, caches MUST not return a cached copy
 - **only-if-cached**: user wants item if cached, otherwise an error
 - **must-revalidate**: server says response may be cached, but must be revalidated once it is stale
 - **proxy-revalidate**: does not apply to user's browser cache

How is Caching Done?

Caching Decisions

- check to see if requested object is in cache
- check if client headers allow item to be returned
- perform **cache coherence** checks
- perform **cache replacement** if needed

Cache Coherency

- cache must ensure that what is in the cache is consistent with what the server stores
- validating
 - [If-Modified-Since](#): using Date
 - [If-Match](#): using ETag
- when should the server validate?
 - use a TTL to indicate how much longer the cached response will be valid
 - based on [Expires](#), [max-age](#) directive, or heuristic that examines last modification time and frequency of requests
- see [Squid Cache](#) FAQ for details on Squid coherence algorithm

Cache Replacement Algorithm

- many important factors
- **access history**: keep objects that are frequently accessed
- **expiration time**: remove objects that will expire soon
- **time since last modification**: keep objects that do not change frequently
- **cost of fetching the resource**: keep in cache if it was expensive to fetch
- **cost of storing the resource**: removing large objects frees a lot of space, but they are expensive to retain

Common Cache Replacement Algorithms

- **Least Recently Used (LRU)**
 - mark objects with time of last access
 - evict object that is least recently accessed
 - old and proven in many areas of CS
 - studies show it is not the best for web caching
- **Least Frequently Used (LFU)**
 - mark objects with how frequently accessed in a given period of time
 - evict object that is least frequently used
- **Size of Object (SIZE):** evict largest object
- **Hyper-G:** first LFU, then LRU, then largest
- **Greedy-Dual Size**
 - compute a utility value for each object
 - evict object with lowest utility
 - utility uses cost of fetching, size, age


Cache Replacement Lessons

- *memory is cheap*: create a really large cache
- *lots of traffic isn't cacheable*
- *most algorithms are good enough*
- Squid uses LRU, Greedy-Dual Size, LFU with Dynamic Aging

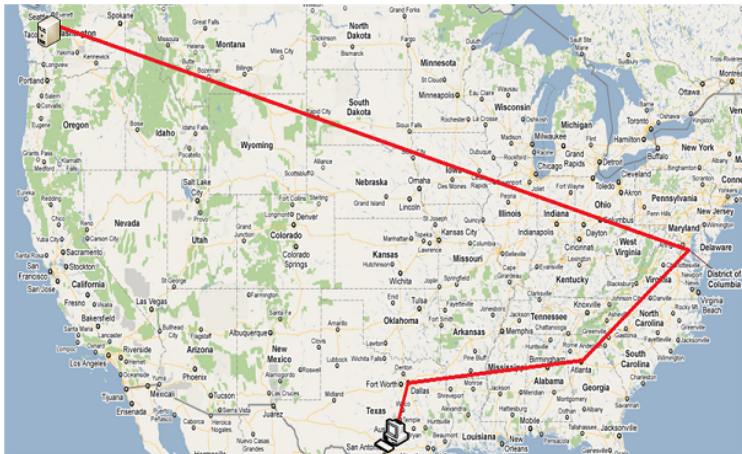
Content Delivery Networks

Web Site Performance

- Steve Souders
(Yahoo, Google):
80% of web page
download time is
spent fetching
embedded images and
scripts



Latency



- want to avoid latency caused by long paths

Content Delivery Network

- replicate content at many caches, typically at the edge of the network
- use domain name of the CDN in your web pages
- client requests routed to a “nearby” server, generally through DNS, reducing loss and delay
- started with Akamai: [IEEE Internet Computing paper](#)
- see [Amazon CloudFront](#)