

Designing Document Database Models

Daniel Zappala

CS 360 Internet Programming
Brigham Young University

Document Databases

Document Databases

- store key/value pairs, and the value is generally a JSON document

```
1  {
2    '_id': <ObjectId>,
3    'version': '2.7',
4    'user_id': 'KWCH-G25',
5    'api_id': 'KL1S-Z4L',
6    'api_data': {...},
7    'name': {'given': 'Carmelo', 'full': 'Carmelo Zappal\xe0', 'family': 'Zappal\xe0'},
8    'gender': 'Male',
9    'living': 'False',
10   'facts': {'death': {'date': '16 March 1983', 'year': '1983', 'place': 'Burbank,
                        Los Angeles, California, United States', 'day': '16', 'month': '3'}, 'birth': {'date': '6 October 1886', 'year': '1886', 'place': 'Linguaglossa,
                        Catania, Sicily, Italy', 'day': '6', 'month': '10'}}
11   ...
```

Data Models

- embedded
 - store related information in a single document
 - useful for one-to-one relationships
 - useful for one-to-many relationships when the child documents always appear with the parent, e.g. comments on a blog
- normalized
 - store related information across separate documents
 - use references to create relationships
 - useful for one-to-many relationships when the children are accessed separately, e.g. want to list a particular user's comments on all blog entries
 - useful for many-to-many relationships
 - useful for very large hierarchical datasets, where the size of a document would exceed database limits

PyMongo Example

```
1  from pymongo import MongoClient
2  client = MongoClient()
3
4  // get a database
5  db = client.blog
6
7  // get a collection
8  posts = db.posts
9
10 // create a post
11 import datetime
12 post = {"author": "Mike",
13         "text": "My first blog post!",
14         "tags": ["mongodb", "python", "pymongo"],
15         "date": datetime.datetime.utcnow()}
16
17 // insert
18 posts.insert(post)
19
20 // get all posts
21 posts.find()
22
23 // get posts by Mike
24 posts.find({'author': 'Mike'})
```

Why Use a Document Database?

- scale
 - cloud computing provides cheap storage, ability to easily add more servers
 - but data must be spread across multiple servers, difficult to use SQL joins in across distributed set of tables
- unstructured data
 - social media, multimedia
 - SQL is best suited for structured data
- agile development
 - database scheme must rapidly evolve
 - easier to do when all items in a table don't need to maintain consistent structure

Document Databases are Good At...

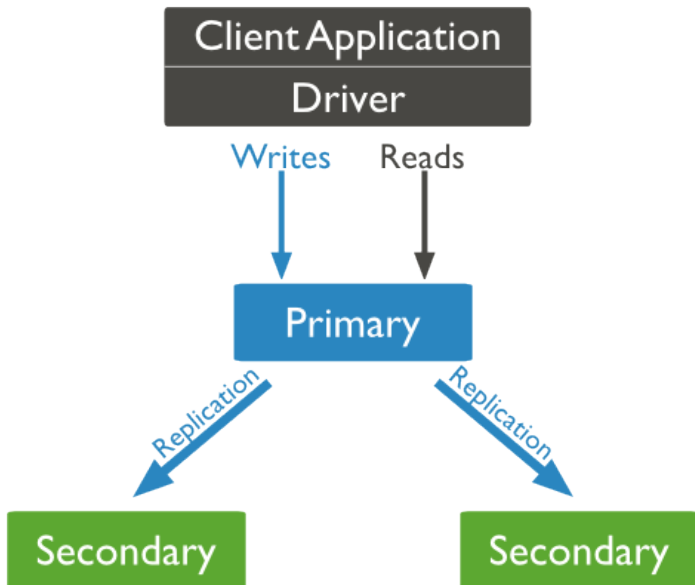
- independent documents
 - fast read performance
 - easy to distribute across servers
- easy application logic
 - object-oriented syntax translates directly into document storage
- unstructured data
 - documents can store whatever keys and values the app requires
 - migrations not needed in advance, can adapt on the fly

Consistency


- CAP Theorem
 - a database can only have two of Consistency, Availability, and Partition tolerance
 - under Partition, a system cannot maintain both Availability and Consistency
- strongly consistent
 - favor consistency over availability
 - relational database
 - MongoDB
- eventually consistent
 - favor availability over consistency
 - Amazon Dynamo, CouchDB, Riak
- easier to provide even load distribution and multi-data center support in eventually consistent systems
- easier to write code for strongly consistent systems, applicable to a wider set of problems

How Mongo Works


Replication




Heartbeats



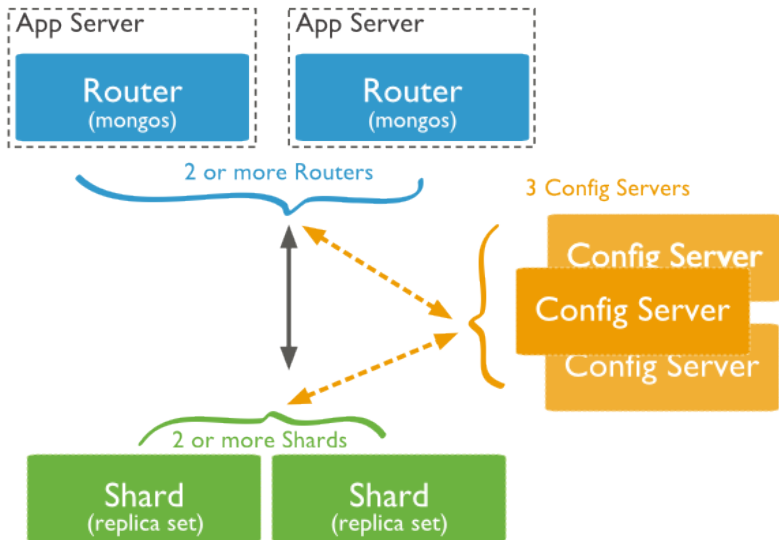
Failover



Sharding




Sharding Configuration



How Riak Works


Storage

- all nodes in a Riak cluster are equivalent; there is no master
- data is distributed across nodes using consistent hashing
 - stores key/value pairs in a bucket namespace (roughly a table)
 - hashes bucket/key combination into a 160-bit integer space
 - each node responsible for portions of the space
- 4-node cluster with 32 partitions:



Replication

- writes specify a replication number W , specifying amount of replication
- stored in that number of consecutive locations
- nearby virtual locations map to physical locations that may be geographically distinct



Eventual Consistency

- reads specify a level of consistency, R
 - specifies how many replicas must return results for a successful read
 - send a request to all nodes where data is stored, return when r have responded
- $W + R > N$ ensures strict consistency
- $W + R \leq N$ provides eventual consistency for better availability or lower latency
- a write may occur during a partition
 - reads will return different values in each partition
 - when the partition is healed, write ordering will eventually result in consistent data

Example

Example

- ▶ Sample FamilyTree App model
- uses ▶ PyMongo