

Thread Synchronization

CS 360 Internet Programming


Daniel Zappala

CS 360 Internet Programming
Brigham Young University


Semaphores

- **semaphore** is a shared variable maintained by OS
 - contains an integer and a queue
 - value initialized ≥ 0
- **wait(s)**: wait for a signal on semaphore s
 - decrements semaphore, blocks if value < 0
 - if blocked, process put on the queue, suspends until signal is sent
- **signal(s)**: transmit a signal to semaphore s
 - increments semaphore
 - if value ≤ 0 then unblock someone
- **wait()** and **signal()** are atomic operations and cannot be interrupted

wait()



signal()



Types of Sempahores

- binary semaphore
 - only one process at a time may be in the critical section
- counting semaphore
 - a fixed number of processes > 0 may be in the critical section
- OS determines whether processes are released from queue in FIFO order or otherwise; usually FIFO in order to prevent starvation

Using Semaphores

```
1 semaphore s = 1;
2
3 void thread(int i) {
4     while (true) {
5         wait(s);
6         /* critical section */
7         signal(s);
8         /* remainder */
9     }
10 }
```

- semaphore protects critical section
- can set s to > 1 to let more than one process in the critical section
 - $s \geq 0$: number that can enter
 - $s < 0$: number that are waiting

POSIX Semaphores

POSIX Semaphores

```
1 #include <semaphore.h>
2
3 int sem_init(sem_t *sem, int pshared, unsigned int value);
4 int sem_wait(sem_t * sem);
5 int sem_trywait(sem_t * sem);
6 int sem_post(sem_t * sem);
```

- `sem_init()`: sets initial value of semaphore; `pshared = 0` indicates semaphore is local to the process
- `sem_wait()`: suspends process until semaphore is > 0 , then decrements semaphore
- `sem_trywait()`: returns `EAGAIN` if semaphore count is $= 0$
- `sem_post()`: increments semaphore, may cause another thread to wake from `sem_wait()`

Example Code

- see example code `semaphore.cc`

► [GitHub](#)

Producer Consumer

Producer Consumer Problem

- one or more producers are generating data and placing them in a buffer
- one or more consumers are taking items out of the buffer
- only one producer or consumer may access the buffer at any time

Producer Consumer

```
1 vector buffer;  
2 append(item) {  
3     buffer.append(item);  
4 }  
5 take() {  
6     return buffer.remove();  
7 }
```

producer:

```
1 while (true) {  
2     item = produce();  
3     append(item);  
4 }
```

consumer:

```
1 while (True) {  
2     item = take();  
3     consume(item);  
4 }
```

Producer Consumer with Infinite Buffer

```
1 sem_t s, n;  
2 sem_init(&s, 0, 1);  
3 sem_init(&n, 0, 0);
```

producer:

```
1 while (True) {  
2     item = produce();  
3     sem_wait(&s);  
4     append(item);  
5     sem_post(&s);  
6     sem_post(&n);  
7 }
```

consumer:

```
1 while (True) {  
2     sem_wait(&n);  
3     sem_wait(&s);  
4     item = take();  
5     sem_post(&s);  
6     consume(item);  
7 }
```

Looking at the Code ...

- ① *What is the purpose of semaphore s ?*
- ② *What is the purpose of semaphore n ?*
- ③ *Why is semaphore s initialized to 1 but semaphore n is initialized to 0?*
- ④ *Why can the producer signal n every time an item is added to the buffer ?*
- ⑤ *Can the producer swap the signals for n and s ?*
- ⑥ *Can the consumer swap the waits for n and s ?*

Important Insights

- two purposes for semaphores
 - **mutual exclusion**: semaphore s controls access to critical section
 - **signalling**: semaphore n coordinates when the buffer is empty: consumer waits if buffer is empty, producer signals when buffer becomes non-empty
- avoid race conditions
 - *item* keeps a local copy of the data protected by the semaphore so that it can be accessed later
 - reduces amount of processing inside the critical section

Important Insights

- n : semaphore value is number of items in buffer
 - if $n == 0$, consumer must wait
 - can swap `sem_post(&n);` and `sem_post(&s);` in producer and be OK
 - can't swap `sem_wait(&n);` and `sem_wait(&s);` in consumer: otherwise consumer enters and then waits and deadlocks the producer!
- ordering of semaphore operations is important

Producer Consumer with Finite Buffer

```
1 sem_t s, n, e;  
2 sem_init(&s, 0, 1);  
3 sem_init(&n, 0, 0);  
4 sem_init(&e, 0, BUFFER_SIZE);
```

producer:

```
1 while (True) {  
2     produce();  
3     sem_wait(&e);  
4     sem_wait(&s);  
5     append();  
6     sem_post(&s);  
7     sem_post(&n);  
8 }
```

consumer:

```
1 while (True) {  
2     sem_wait(&n);  
3     sem_wait(&s);  
4     take();  
5     sem_post(&s);  
6     sem_post(&e);  
7     consume();  
8 }
```

Looking at the Code ...

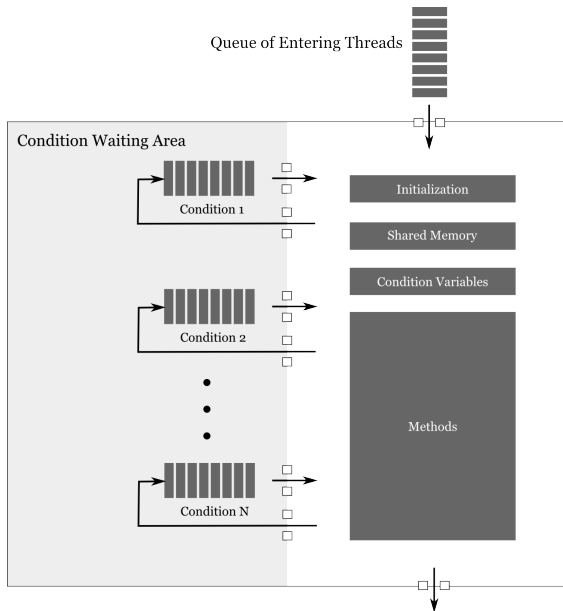
- 1 *What is the difference between semaphore e and semaphore n ?*

Monitors

Monitor

- difficult to get semaphores right
 - match wait and signal
 - put in right order
 - scattered throughout code
- **monitor**: programming language construct
 - equivalent functionality
 - easier to control
 - mutual exclusion constraints can be checked by the compiler
 - used in versions of Pascal, Modula, Mesa
 - Java also has a Monitor object but compliance cannot be checked at compile time

Hoare Monitor



Hoare Monitor

- monitor can only be entered through methods
- shared memory can only be accessed by methods
- only one process or thread in monitor at any time
- may suspend and wait on a condition variable
- like object-oriented programming with mutual exclusion added in

Hoare Synchronization

- `cwait(c)`: suspend on condition `c`
- `csignal(c)`: wake up one thread waiting for condition `c`
 - do nothing if no threads waiting (signal is lost)
 - different from semaphore (number of signals represented in semaphore value)

Producer Consumer with a Hoare Monitor

```
1 vector buffer;  
2 condition notfull, notempty;
```

```
1 append(item) {  
2   if buffer.full()  
3     cwait(notfull);  
4   buffer.append(item);  
5   csignal(notempty);  
6 }
```

```
1 take() {  
2   if buffer.empty();  
3     cwait(notempty);  
4   item = buffer.remove();  
5   csignal(notfull);  
6   return item;
```

Producer Consumer with a Hoare Monitor

producer:

```
1 while (True) {  
2   item = produce();  
3   append(item);  
4 }
```

consume:

```
1 while (True) {  
2   item = take();  
3   consume(item);  
4 }
```

- advantages
 - moves all synchronization code into the monitor
 - monitor handles mutual exclusion
 - programmer handles synchronization (buffer full or empty)
 - synchronization is confined to monitor, so it is easier to check for correctness
 - write a correct monitor, any thread can use it

Lampson and Redell Monitor

- Hoare monitor requires that signaled thread must run immediately
 - thread that calls `csignal()` must exit the monitor or be suspended
 - for example, when `notempty` condition signaled, thread waiting must be activated immediately or else the condition may no longer be true when it is activated
 - usually restrict `csignal()` to be the last instruction in a method (Concurrent Pascal)
- Lampson and Redell
 - replace `csignal()` with `cnotify()`
 - `cnotify(x)` signals the condition variable, but thread may continue
 - thread at head of condition queue will run at some future time
 - must recheck the condition!
 - used in Mesa, Modula-3

Producer Consumer with a Lampson Redell Monitor

```
1 vector buffer;  
2 condition notfull, notempty;
```

```
1 append() {  
2   while buffer.full()  
3     cwait(notfull);  
4   buffer.append(item);  
5   cnotify(notempty);  
6 }
```

```
1 take() {  
2   while buffer.empty()  
3     cwait(notempty);  
4   item = buffer.remove();  
5   cnotify(notfull);  
6   return item;  
7 }
```

Lampson Redell Advantages

- allows processes in waiting queue to awaken periodically and reenter monitor, recheck condition
 - prevents starvation
- can also add `cbroadcast(x)`: wake up all processes waiting for condition
 - for example, append variable block of data, consumer consumes variable amount
 - for example, memory manager that frees k bytes, wake all to see who can go with k more bytes
- less prone to error
 - process always checks condition before doing work

What Can You Do?

- emulate a Lampson Redell Monitor with semaphores
 - create a class with private data only
 - use the same semaphore to protect all class methods
 - use semaphores to replace `cwait()` and `cnotify()`
 - use while loops to recheck conditions
- take your semaphores and move them inside the method call instead of outside of it