

# Socket Programming

## CS 360 Internet Programming

Daniel Zappala

Brigham Young University  
Computer Science Department

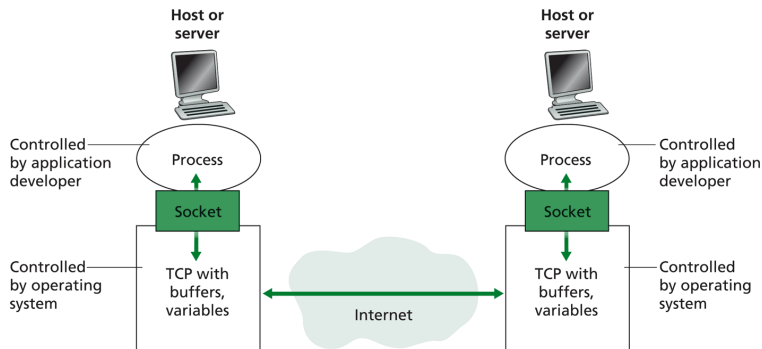
# Sockets, Addresses, Ports

# Clients and Servers



- clients request a service from a server using a protocol
- need an API for sending and receiving data
- need an abstraction for the Internet
  - a reliable connection: TCP
  - an unreliable message service: UDP
- these are provided by the **BSD socket API**

# Processes and Sockets



- inter-process communication uses messages sent on sockets
- socket API defines
  - how to open, close, read, and write to socket
  - which transport protocol to use
  - various communication parameters


# Socket API

- BSD Socket API - the dominant socket API on Linux, BSD, Windows
- Use the man pages!
  - shows C syntax, with include files
  - gives a description of how the system call works
  - lists return values
  - lists errors
  - lists other relevant man pages
- *man man*
- *man socket*
- the socket API is in section 2 of the manual, so you will need to specify the section if the same command exists in an earlier section: *man 2 bind*

# Addresses and Ports

- to talk to a process on another machine, you need to identify it
  - **IP address**: identifies the machine
  - **port number**: identifies the socket the process is using

# IP Addresses



- IPv4 address identifies an interface/link on a host or router
  - 32 bits
  - dotted-decimal notation: each part is 8 bits

# Ports

- identifier for a socket on which a process is listening
  - 16 bits
  - a process may listen on more than one socket; each must be on a separate port
- Operating systems (such as Linux) designate some ports as *privileged*, meaning only the superuser can listen on them (typically ports less than 1024)
- To help find common servers, the IANA designates *well-known* ports for many protocols
  - HTTP: 80
  - SMTP: 25
  - SSH: 22
  - NTP: 123



**ClientAPI**

# Client API

- ① `socket()`: create a socket endpoint
- ② `connect()`: connect to a server
- ③ `send()`: send data
- ④ `recv()`: receive data
- ⑤ `close()`: close the socket

# Creating a Socket

---

```
1 int socket(int domain, int type, int protocol);
2
3 domain = PF_INET
4 type = SOCK_STREAM for TCP, SOCK_DGRAM for UDP
5 protocol = 0
```

---

- on success returns a socket descriptor

# Connecting to a Server

---

```
1  int  connect(int sockfd, const struct sockaddr *serv_addr ,
2  socklen_t addrlen);
3
4  sockfd = socket you created
5  serv_addr = pointer to socket address structure
6  socklen_t = length of socket address
```

---

- connects to a server
- uses the socket address structure to pass an IP address and port
- on success returns zero
- on error returns -1 and sets errno

# Sending and Receiving Data

- usually the client sends a request to the server and the server sends a reply
- socket operations are similar to reading from and writing to a file
  - a socket descriptor acts like a file handle
  - sending  $\sim$  writing
  - receiving  $\sim$  reading

# Send and Receive Syntax

---

```
1 ssize_t send(int s, const void *buf, size_t len, int flags);  
2 ssize_t recv(int s, void *buf, size_t len, int flags);
```

---

- `s` = socket
- `buf` = pointer to buffer
- `len` = size of buffer in bytes
- on success returns the number of bytes actually sent or received
  - if it is less than what you expected, then you must repeat the system call until all data is sent or received
- `recv()` will return 0 if the socket has been closed
- on error returns -1 and sets `errno`
- see man pages for more advanced options

# Closing a Socket

---

```
1 #include <unistd.h>
2
3 int close(int fd);
```

---

- `fd` = socket
- on success returns zero
- on error returns `-1` and sets `errno`
- releases the socket file descriptor so it can be re-used!

# Example

- *See simple echo client on github.*

► [GitHub](#)



# Socket Addresses

# Socket Addresses

- to connect a socket, you must supply a socket address structure that includes the IP address and port of the server you are connecting to
  - ① use DNS to convert host name to IP address
  - ② initialize a socket address structure

# Generic Socket Address Structure

---

```
1 struct sockaddr {  
2     sa_family_t    sa_family;    // address family  
3     char           sa_data[14];  // address  
4 }
```

---

- used to represent a generic connection between processes
- potentially provides access to many different addressing standards
- address family can be
  - [AF\\_UNIX](#) : local UNIX socket
  - [AF\\_INET](#) : Internet socket

# IPv4 Socket Address Structure

---

```
1 struct sockaddr_in {
2     sa_family_t    sin_family;    // address family
3     u_int16_t      sin_port;      // port
4     struct in_addr  sin_addr;      // Internet address
5     char           sin_zero[8];    // unused
6 }
```

---

- can cast IPv4 socket into a generic socket
- **port** is 2 bytes, **address** is 4 bytes, **zero** is 8 bytes = 14 bytes
- IPv4 address structure:

---

```
1 struct in_addr {
2     u_int32_t      s_addr;    // address
3 };
```

---

# Network Byte Order

- you must store the port and address in network byte order (most significant byte sent first)
- provides interoperability among Internet hosts
- use `htons()` for the port and `inet_aton`, `inet_addr()` or `inet_makeaddr()` for the address

---

```
1 struct sockaddr_in server;  
2 server.sin_port = htons(80);  
3 if (!inet_aton(ipaddress, &server.sin_addr))  
4     printf("inet_addr() conversion error\n");
```

---

# Example

- *See simple echo client on GitHub.*

► [GitHub](#)

# Server API

# Server API

- 1 create a `socket`
- 2 `bind`: associate the socket with an address and port
- 3 `listen`: convert socket to listen for incoming connections
- 4 `accept`: accept an incoming client connection
- 5 `send`: send data
- 6 `recv`: receive data



# Binding the Socket

---

```
1  int bind(int sockfd, const struct sockaddr *my_addr,
2          socklen_t addrlen);
3
4  sockfd = socket you created
5  my_addr = pointer to socket address structure
6  addrlen = length of socket address
```

---

- associates an address with a socket
- uses the socket address structure to pass an IP address and port
- on success returns zero, -1 and errno otherwise

# Listening on the Socket

---

```
1  int listen(int sockfd, int backlog);  
2  
3  sockfd = socket you created and bound  
4  backlog = maximum waiting connections
```

---

- converts a socket to a passive socket (one that accepts connections rather than connecting)
- backlog is the maximum number of waiting connections the kernel should hold in a queue
- maximum value in Linux is 128
- on success returns zero, -1 and errno otherwise

# Accepting a Connection

---

```
1  int accept(int sockfd, struct sockaddr *addr,  
2             socklen_t *addrlen);  
3  
4  sockfd = socket you created and bound  
5  addr = pointer to empty socket address structure  
6  addrlen = length of empty socket address structure
```

---

- accept a connection from a client; gets the next connection waiting in the queue
- if there are no pending connections, the process sleeps assuming this is a blocking socket (this is the default)
- on success returns a new socket descriptor
- on success, `accept()` fills in the address of the client

# Example

- *See simple echo server on GitHub.*

► [GitHub](#)

# Parsing Messages

# Example

- *Try simple echo client with slow echo server in the socket example code.*

► [GitHub](#)

# Parsing Messages

- an application protocol reads and write *messages*, but TCP and the socket API use a *byte stream*
  - **when a client calls `recv()` there is no guarantee that it will get an entire message**
  - the application has to designate where a message starts and ends and then parse the byte stream looking for messages
- two options for reading a message
  - ① variable length: **read until a sentinel** (end-of-message marker)
  - ② known length: read a length field and then **read the listed number of bytes**
- **you will need to write a `read-until-sentinel()` and `read-fixed-length()` methods for your socket programming labs**

# Sending and Receiving Properly

- results of `send()` or `recv()` call
  - ① less than zero bytes, `errno == EINTR`
    - the system call was interrupted – try again
  - ② less than zero bytes, any other error
    - fatal error – try to recover gracefully
  - ③ zero bytes
    - the socket is closed – try to recover gracefully
  - ④ positive bytes
    - the return value is the number of bytes sent or received
- **use a `send()` or `recv()` loop!**



# Example

- *See echo server and client on GitHub.*

► [GitHub](#)