

# Performance Measurement and Workload Models

Daniel Zappala

CS 360 Internet Programming  
Brigham Young University

# Performance Measurement

# Motivation: Content Creators

- goal: maximize profit
- how do people use my web site?
  - what led people to buy this book? recommendation list? ad on another site?
  - what led people to *not* buy this book? negative reviews? ordering process too complicated?
- how many people use my web site?
  - ad revenue
  - where do they come from?
  - where are *not* coming from?

# Motivation: Web Hosting and ISPs

- goal: maximize profit and performance
- accounting
  - bill people for the bandwidth they use
- system performance
  - replicate content to balance load
  - evaluation of a server to find and examine bottlenecks
  - effect of server parameters, e.g. number of active threads
  - measure the effectiveness of caching: number of hits, bandwidth saved

# Motivation: Network Researchers

- goal: understand the network and how well it is working
- understanding
  - determine how many hosts have deployed a new protocol or feature (e.g. how many caches use Etag for validation?)
  - measure and model Internet traffic
- protocol development: improve performance
  - led to HTTP/1.1 persistent connections
  - major role in determining TCP performance
  - evaluate cache replacement, validation, and prefetching algorithms

# Types of Measurement

# Passive Measurement

- examine log files
- place a monitor “inside” the network
  - promiscuous mode on Ethernet means you get all packets
  - modify host to log all IP packets to a very large disk
- advantages
  - can capture all HTTP and TCP behavior
  - can examine a detailed timeline of all events
  - non-invasive for the web server and clients
- limitations
  - hard to capture all packets at line speed
  - limited to examining what happened, not what *may* happen

# Active Measurement

- generate user requests on your own to examine web server and user performance
- locating user agents
  - wide variety of connection speeds
  - wide variety of network locations
  - what distributions?
- generating requests
  - which web pages?
  - how many requests per client?
  - use distributions collected from past experiments
- collecting data
  - DNS time
  - TCP connection time
  - HTTP request time



# Simulation

- model a system and examine its behavior in a simulator
- requires modeling
  - the Internet (connectivity, bandwidth, delay)
  - Internet traffic
- complexity depends on level of simulation
  - packet-level simulation: treat each packet as an event, which must be handled at each router and host in the network
  - session-level simulation: treat each file transfer as an event, which must be handled only at the hosts in the network

# Simulation vs Measurement

- advantages of simulation
  - network can be tightly controlled delay, bandwidth, competing traffic
  - can often simulate at a larger scale than a measurement study
  - experiments in a simulated network run in simulated time (much faster than real time)
  - can often vary more factors for a more complete study
- advantages of measurement
  - you learn different things when you actually have to build it
  - more convincing results: it works in the “real world”
  - can use real users and real traffic

# Deciding Between Simulation and Measurement

- can you implement your solution?
  - time allocated to design, spec, testing, debugging
  - is there an existing framework or library that would make your job easier?
- can you deploy your solution?
  - do you have access to enough computers or friends?
  - what operating system and libraries will you use?
  - can you use PlanetLab?
- how many factors do you want to study?

# Methodology

# Experimental Methodology

- important components of an experimental methodology for network performance measurement
  - ① topology
  - ② workload
  - ③ data collection
  - ④ metrics

# Topology

- what does the network look like?
  - LAN
  - campus
  - PlanetLab
  - Internet
- simulated networks must be generated
  - how many nodes?
  - how many links?
  - link and node characteristics - bandwidth, delay, queue size
  - network characteristics - organizations, ISPs

# Workload

- what load will you put on your solution and therefore on the network?
- *example: web server*
  - how often do clients make a request to the web server?
    - constant arrival rate
    - exponential arrival rate
  - how many requests does a client make and what is the time between requests?
  - what size are the objects that are available on the web server?
  - what is the popularity of each of the objects?
  - what is the available bandwidth to each client?

# Data Collection

- collect a log of important events in your system
- **important: collect data, not metrics**
  - if you collect metrics, you will invariably find that you need to change a metric or add another one and now you need to re-run all of your experiments
  - collecting data takes more room but if you generalize it then you can compute any metrics you want later on
- *example: web server*
  - thread time threadID start
  - thread time threadID stop
  - queue time size



# Metrics

- how do you know whether your solution is solving the problem?
- how do you know it is doing it better than other solutions?
- *hint*: look at previous solutions and see what metrics they use
- *example: web server*
  - response time
  - client bandwidth
  - percentage of time each thread is busy
  - queue size over time
  - average, standard deviation, percentiles, distributions, etc.
- note: knowing whether a difference is significant requires a good statistics background

# Workload

# Two Approaches

- trace-driven workload
  - use a log of actual requests to a server
  - reproduces an actual workload
  - avoids statistical analysis
  - no flexibility
  - may not be representative of other situations
  - mingles load with performance (the rate at which a user can make requests depends on how fast it can download a page)
  - hard to collect, requires extensive storage
- stress testing
  - send requests as fast as possible to induce high load
  - can quickly identify performance bottlenecks
  - not realistic

# Synthetic Workloads

- characterize the load placed on a system using a mathematical/statistical model
- want it to be repeatable, realistic, parameterized
- advantages
  - derived from a model that can be analyzed, validated, and compared to other models
  - can capture general properties of real traffic without being limited to a particular trace
  - can explore performance in a controlled setting by adjusting the parameters of the model

# Probability Distributions

# Modeling Workload

- *example: web server log - how can we characterize the size of a response message?*
- mean can be skewed by high variance
  - response messages: 4100, 4700, 4200, 20,000, 4000 bytes
  - mean size = 7400 bytes
- median doesn't represent variance
  - median size = 4200 bytes
  - response messages: 4100, 4700, 4200, 4800, 4000 bytes has same median
- variance, standard deviation provide more information but they are also relatively simple
- need probability distributions to provide a more complete picture

# Probability Density Function (pdf)


- describes the likelihood for a continuous random variable to occur over an interval
- defined by a function  $f(x)$  such that
  - $f(x) \geq 0$  for all values of  $x$
  - the total area under the graph is 1

$$\int_{-\infty}^{\infty} f(x) dx = 1 \quad (1)$$

- probability of the interval  $(a, b]$  is given by  $\int_a^b f(x) dx$  for any two numbers  $a$  and  $b$

# Cumulative Distribution Function (CDF)

- describes the likelihood that a continuous random variable will have a value less than or equal to  $x$
- $F(x) = P(X \leq x)$







# Common Distributions



# Exponential Distribution

- $F(x) = e^{-\lambda x}$
- frequently used to model the time interval between successive random events that happen at a constant average rate
- mean  $\mathbf{E}[X] = \frac{1}{\lambda}$



# Pareto Distribution

- frequently used to model resource sizes
  - Italian economist Vilfredo Pareto (1848 - 1923)
  - 1906 observation: *20% of the population own 80% of the property*
- $P(X > x) = (\frac{x}{b})^{-a}, x \geq b > 0, b = \text{minimum possible value}$ 
  - also called a “heavy-tailed” distribution
  - Wikipedia uses  $a = k, b = 1$  in graphs below



# Generating Random Variates

- want to be able to generate a random workload that follows a particular distribution
  - *generate the arrival of the next web request using an exponential distribution*
- exponential function
  - $x = -\frac{1}{\lambda} \ln(1 - U)$
  - $U$  is a uniform random variable  $[0,1]$
  - choose a random number between 0 and 1, and calculate the corresponding value of  $x$
- Pareto
  - $x = \frac{b}{(1-U)^{\frac{1}{a}}}$

# Discrete Probability

# Discrete Probability

- probability mass function (pmf)
  - probability that a discrete random variable is exactly equal to some value
  - rank resources from highest to lowest probability and plot probability for each rank
- typically use Zipf's Law to model resource popularity

# Zipf's Law

- in a corpus of natural language, the frequency of a word is inversely proportional to its rank
  - Harvard linguist George Kingsley Zipf
  - *example: in the Brown Corpus, "the" accounts for 7% of all words, "of" accounts for about 3.5% of all words, etc*
- $P(r) = kr^{-1}$ , where  $r$  is the rank of an object and  $k$  is constant such that all probabilities sum to 1
  - more generally,  $P(r) = kr^{-c}$  for some constant  $c$
  - heavy-tailed
- similar to Pareto
  - Zipf-like distributions can be viewed as Pareto distributions if you exchange some variables

# Modeling Workload



# Determining the Right Distribution

- ① determine workload parameters
- ② measure some real traffic
- ③ analyze measurement data to construct a statistical model of the workload
- ④ validate the model against additional real workloads

# Generating a Web Workload

- ① generate session (user) arrivals using an exponential distribution
- ② for each session, generate the number of pages and the number of objects per page using Pareto distributions
- ③ for object size, use a Pareto distribution
- ④ for object popularity use a Zipf distribution